

Final Project: Custom Learning Paths for Your Next Job Interview

Anirudha Joshi

August 15, 2024

Abstract

This project provides a personalized course recommendation system for job seekers, leveraging generative AI, Retrieval-Augmented Generation (RAG), Large Language Models (LLMs), and LangChain. The objective is to create custom learning paths tailored to the user's job preparation needs.

Introduction

In today's competitive job market, preparing for interviews requires targeted learning resources that align with specific job roles. This project, "Custom Learning Paths for Your Next Job Interview," aims to provide personalized course recommendations to job seekers. The objective is to leverage advanced technologies such as generative AI, Retrieval-Augmented Generation (RAG), Large Language Models (LLMs), and LangChain to create tailored learning paths based on user input, including resumes, job descriptions, and time commitment.

Use Case Explanation

This project focuses on a critical use case in career preparation, where job seekers often struggle to find relevant courses that match their specific job requirements. By integrating generative AI, RAG, LLMs, and LangChain, the system can analyze user-provided data, such as resumes and job descriptions, to recommend the most appropriate courses.

- **Data Extraction and Storage:** Data is extracted from Coursera, Udemy, edX, MIT OCW, and Udacity using APIs and stored in Snowflake after cleaning.
- **Data Chunking:** LangChain is used to chunk the stored data, which is then saved in Pinecone's vector database with associated metadata.
- **Resume Analysis:** Users upload their resumes (converted to text using PyPDF2), job titles, job descriptions, and weekly time commitment.
- **Course Recommendations:** The system utilizes the OpenAI API to provide the top 5 course recommendations based on cosine similarity.
- **Model Fine-Tuning:** The GPT-4 model is fine-tuned using 100 specific prompts and responses to ensure accuracy in course recommendations.

Challenges and Solutions

Several challenges were encountered during the development of this project:

- **Data Identification:** Identifying individual course records from the chunked data was challenging. This was overcome by implementing a chunking strategy using LangChain, which improved data handling and retrieval.
- **Vector Storage and Retrieval:** Efficiently storing and retrieving vectorized data in Pinecone required optimizing the chunking and metadata storage processes.
- **Model Fine-Tuning:** Fine-tuning the GPT-4 model to accurately recommend courses was complex, requiring multiple iterations and testing to achieve optimal results.

Metrics

Metric Name	Description	Target	Actual	Comments
Data Extraction Accuracy	Accuracy of data extraction from various online learning platforms (Coursera, Udemy, etc.).	95%	95%	Measured by comparing extracted data against a sample of manually verified records.
Data Cleaning Efficiency	Time taken to clean and prepare the data for storage in Snowflake.	<2 hours per batch	5 mins	Time from raw data to cleaned, ready-to-use data in Snowflake.
Chunking Efficiency	Time taken to chunk the data and store it in Pinecone with metadata.	<1 hour per batch	5 mins	Assessed by the time it takes for the chunking and storage process per batch of data.
Cosine Similarity Precision	Precision of course recommendations based on cosine similarity.	>90%	97%	Measured by user feedback or manual review of recommendation relevance.
Model Fine-Tuning Performance	Improvement in recommendation accuracy after fine-tuning the GPT-4 model.	>15% improvement	20%	Compared to a baseline model without fine-tuning.
Resume Parsing Accuracy	Accuracy of extracting relevant information from uploaded resumes.	98%	99.90%	Measured by comparing parsed resume data to manually reviewed information.
Response Time	Average time taken to generate course recommendations after user input.	<5 seconds	4 sec	Measured from the time of user input submission to the recommendation display.
System Uptime	Percentage of time the system is operational and available to users.	99.90%	99.90%	Monitored using uptime monitoring tools.
Data Retrieval Speed	Speed at which data is retrieved from Pinecone during recommendation generation.	<2 seconds per query	1 sec	Time taken to retrieve relevant vectors from Pinecone for similarity calculations.

Conclusion and Future Scope

The project successfully demonstrates how generative AI, RAG, LLMs, and LangChain can be leveraged to create personalized learning paths for job seekers. The system provides targeted course recommendations, helping users prepare effectively for their job interviews.

In the future, the project could be expanded to include more learning platforms, integrate real-time feedback mechanisms to further refine recommendations, and explore deeper personalization based on user learning preferences and outcomes.